Paying Not to Go to the Gym: Present Bias, Contract Choice, and Welfare in Health Club Membership

Adriana Delgado-Voss, Kweku Mensah, Sophia Bergstrom

March 16, 2026

Abstract

We use health club membership data from three US gyms (1997-2001) to estimate present bias and its welfare consequences. Members who choose flat-rate monthly contracts attend on average 4.17 times per month, well below the break-even of 4.5 visits needed to justify the fee over pay-per-visit pricing. Structural estimation of a beta-delta model yields $\hat{\beta} = 0.712$ and an overconfidence parameter $\hat{\theta} = 0.591$, indicating that members are substantially naive about their future present bias. The average welfare loss from present bias is \$136 per member per year, with 62% attributable to dominated contract choice and 38% to procrastination. Counterfactual analysis shows that a commitment contract eliminating present-bias distortions would raise total surplus by 15% and consumer surplus by 32%.

Contents

I. Introduction

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality

or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

$$U_t = u(c_t) + \beta \sum_{\tau=1}^{T} \delta^\tau u(c_{t+\tau}), \quad \mathbb{E}_t[U_t] = \mathbb{E}\left[u(c_t) + \beta \sum_{\tau=1}^{T} \delta^\tau u(c_{t+\tau}) \,\middle|\, \mathcal{F}_t \right] \tag{1}$$

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak

to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

A distinctive feature of our analysis is the richness of the data. Unlike many previous studies that rely on aggregate proxies, our dataset contains individual-level information on both the outcome variable and the treatment variable, measured with high frequency over an extended time period. This granularity allows us to investigate heterogeneity in treatment effects across individuals, to construct precise event study estimates, and to conduct a range of placebo and falsification tests that would not be feasible with more aggregated data.

Before proceeding, we note several contributions relative to the closest existing work. While prior studies have documented a correlation between the treatment variable and the outcome variable in various contexts, our paper is among the first to provide credibly causal estimates using an instrumental variables strategy. Moreover, our analysis of mechanisms goes beyond what has been possible in earlier studies, shedding light on the channels through which the treatment variable affects the outcome variable.

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of

prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

II. Institutional Background and Data

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality

or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

Table 1: Summary Statistics: Health Club Members

	Mean	SD	p10	p90
Panel A: Contract Choice at Enrollment				
Monthly Contract (1=yes)	0.614	0.487	0.00	1.00
Annual Contract (1=yes)	0.241	0.428	0.00	1.00
Pay-Per-Visit (1=yes)	0.145	0.352	0.00	1.00
Panel B: Attendance and Tenure				
Average Monthly Visits	4.17	4.02	0.50	9.50
Monthly Contract Duration (months)	8.22	9.14	1.00	21.00
Average Monthly Fee Paid ($)	70.14	18.32	49.00	94.00
Cost per Visit ($)	17.80	19.41	6.50	42.00
Panel C: Demographic Characteristics				
Age	33.8	10.2	22.0	49.0
Female (1=yes)	0.481	0.500	0.00	1.00
Distance to Club (miles)	2.14	1.88	0.41	4.82
Observations (members)	7,752	7,752	7,752	7,752
Member-month observations	74,661	74,661	74,661	74,661

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

A distinctive feature of our analysis is the richness of the data. Unlike many previous studies that rely on aggregate proxies, our dataset contains individual-level information on both the outcome variable and the treatment variable, measured with high frequency over an extended time period. This granularity allows us to investigate heterogeneity in treatment effects across individuals, to construct precise event study estimates, and to conduct a range of placebo and falsification tests that would not be feasible with more aggregated data.

Before proceeding, we note several contributions relative to the closest existing work.

While prior studies have documented a correlation between the treatment variable and the outcome variable in various contexts, our paper is among the first to provide credibly causal estimates using an instrumental variables strategy. Moreover, our analysis of mechanisms goes beyond what has been possible in earlier studies, shedding light on the channels through which the treatment variable affects the outcome variable.

A Health Club Setting

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable

bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has

been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

B Contract Menu and Pricing

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable

bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

C Survey Data on Attendance Intentions

Our primary data source is administrative records, which covers 50,000 individuals over the period 1990–2010. This dataset provides detailed information on the outcome variable, the treatment variable, and a rich set of control variables that we use to address potential confounders. The data are available at the individual level and include both repeated cross-sections and a longitudinal component that allows us to track individuals over time. We discuss the advantages and limitations of this dataset relative to alternatives in the literature.

Table **??** presents summary statistics for the key variables used in our analysis. The mean of the outcome variable is consistent with values reported in prior studies, providing confidence that our sample is representative of the population of interest. The standard deviation suggests substantial variation across individuals, which is essential for our identification strategy. We also report the 10th and 90th percentiles to characterize the full distribution and to assess the extent to which our results might be driven by outliers in the tails.

We construct our main sample by merging several administrative and survey datasets. After applying standard sample restrictions and dropping observations with missing values on key variables, our final estimation sample contains 50,000 observations covering 1990–2010. The merge rate across datasets exceeds 95 percent, alleviating concerns about selective attrition at the matching stage. We verify that the merged sample is representative by comparing moments of the key variables to those in the underlying source datasets.

The key independent variable, the treatment variable, exhibits substantial cross-sectional variation. The interquartile range spans a factor of three, providing sufficient variation to identify the effect of interest. We verify that this variation is not driven by outliers or measurement error by conducting a series of diagnostic tests. In particular, we show that the results are robust to winsorizing the treatment variable at the 1st and 99th percentiles, to using alternative measures from different data sources, and to correcting for measurement

error using an errors-in-variables estimator.

We supplement the main dataset with auxiliary sources to construct additional control variables and instruments. Geographic controls come from standard sources in the literature, including distance to the coast, latitude, rainfall, temperature, and terrain ruggedness. Demographic controls are measured at baseline to avoid bad control problems. Institutional controls are drawn from widely used cross-country datasets. We lag all time-varying controls by one period to mitigate simultaneity concerns.

The measurement of the outcome variable deserves discussion. Following the standard approach in the literature, we use the logarithm of the outcome variable as our primary outcome variable. This transformation has three advantages: it reduces the influence of outliers, it allows us to interpret coefficients as approximate percentage changes, and it is consistent with the multiplicative structure of the theoretical model. We verify that our results are qualitatively robust to using the outcome variable in levels or in ranks.

Panel A of Table **??** reports summary statistics for the full sample, while Panels B and C restrict attention to subsamples that will be used in the heterogeneity analysis. The key variables are well-balanced across subsamples, with the exception of the treatment variable, which exhibits the expected variation driven by the quasi-experimental design. We formally test for balance across subsamples using a multivariate F-test and fail to reject the null of joint equality at conventional significance levels.

III. Theoretical Framework

Consider an economy populated by a continuum of agents indexed by $i \in [0, 1]$. Each agent chooses the outcome variable to maximize lifetime utility subject to a budget constraint. The key source of heterogeneity across agents is their exposure to the treatment variable. We solve for the equilibrium of this economy and derive comparative statics with respect to the key parameters. The model generates a clear testable prediction: an exogenous increase in the treatment variable should lead to a positive and significant change in the outcome

variable, with the magnitude depending on structural parameters that we estimate in the empirical analysis.

Let Y_i denote the outcome of interest for unit i. We model Y_i as a function of treatment status D_i, observed covariates X_i, and unobserved heterogeneity ε_i. Under standard regularity conditions, the conditional expectation function takes the form presented in Equation **??**. The parameter of interest is the coefficient on treatment status, which captures the causal effect of the treatment variable on the outcome variable holding constant observed and unobserved confounders. Identification of this parameter requires an exclusion restriction, which we discuss in detail below.

The theoretical framework builds on the canonical model of field. Agents face a discrete choice between J alternatives, each yielding utility $U_{ij} = V_{ij} + \epsilon_{ij}$, where V_{ij} is the deterministic component and ϵ_{ij} captures idiosyncratic preferences. Under standard distributional assumptions on the error term, the choice probabilities take a tractable closed form. The model can be estimated using maximum likelihood or method of moments, and the resulting structural parameters can be used to simulate counterfactual policy scenarios.

We derive the key prediction of the model through a simple comparative statics exercise. An exogenous increase in the treatment variable shifts the equilibrium allocation of the outcome variable through two channels: a direct effect operating through prices, and an indirect effect operating through the reallocation of factors across sectors. The total effect is ambiguous in general, but under the empirically relevant parameter values, we show that the direct effect dominates. Proposition 1 in Appendix A formalizes this result and establishes conditions under which the sign of the total effect is unambiguous.

Under the assumption that the proposed instrument satisfies the exclusion restriction, the model implies that the reduced-form relationship between the proposed instrument and the outcome variable is proportional to the structural parameter of interest. This motivates our instrumental variables strategy described in Section **??**. A key advantage of the IV approach is that it does not require functional form assumptions on the relationship between

the treatment variable and the unobserved confounders. The cost is that we identify a local average treatment effect rather than the population average treatment effect.

$$V^M > V^{PPV} \iff \sum_{t=1}^{T} \delta^t \beta[b - c(\nu_t)] - p^M > \sum_{t=1}^{T} \delta^t \beta[b - c(\nu_t)]\mathbf{1}[\nu_t > 0] - p^{PPV}\nu_t \tag{2}$$

$$\nu_t^* = 1 \iff \beta[b - c_t] \geq p_t, \quad c_t \sim F(\cdot) \tag{3}$$

$$tildebeta = theta + (1 - theta)beta, quadthetain[0, 1] \tag{4}$$

The model delivers additional testable implications beyond the main prediction. First, the treatment effect should be larger for individuals with lower baseline levels of the outcome variable, reflecting diminishing marginal returns. Second, the effect should be persistent over time rather than transitory, because the mechanism operates through permanent changes in the level of the treatment variable. Third, the effect should be concentrated in sectors with greater exposure to the underlying channel. We test each of these auxiliary predictions in Section 6.

We make several simplifying assumptions to maintain tractability. First, we assume that agents have rational expectations and perfect information about the distribution of shocks. Second, we abstract from general equilibrium effects by treating prices as exogenous to individual agents. Third, we assume that the production technology exhibits constant returns to scale. We discuss the robustness of our theoretical predictions to relaxing each of these assumptions in Appendix A, where we also present a more general version of the model.

The equilibrium of the model can be characterized by a system of first-order conditions and market-clearing conditions. In the interior solution, the marginal benefit of increasing the treatment variable equals the marginal cost, yielding an implicit equation for the optimal

level of the treatment variable as a function of exogenous parameters and state variables. We log-linearize this equation around the steady state to obtain the estimating equation that motivates our empirical specification.

Consider an economy populated by a continuum of agents indexed by $i \in [0, 1]$. Each agent chooses the outcome variable to maximize lifetime utility subject to a budget constraint. The key source of heterogeneity across agents is their exposure to the treatment variable. We solve for the equilibrium of this economy and derive comparative statics with respect to the key parameters. The model generates a clear testable prediction: an exogenous increase in the treatment variable should lead to a positive and significant change in the outcome variable, with the magnitude depending on structural parameters that we estimate in the empirical analysis.

Let Y_i denote the outcome of interest for unit i. We model Y_i as a function of treatment status D_i, observed covariates X_i, and unobserved heterogeneity ε_i. Under standard regularity conditions, the conditional expectation function takes the form presented in Equation **??**. The parameter of interest is the coefficient on treatment status, which captures the causal effect of the treatment variable on the outcome variable holding constant observed and unobserved confounders. Identification of this parameter requires an exclusion restriction, which we discuss in detail below.

The theoretical framework builds on the canonical model of field. Agents face a discrete choice between J alternatives, each yielding utility $U_{ij} = V_{ij} + \epsilon_{ij}$, where V_{ij} is the deterministic component and ϵ_{ij} captures idiosyncratic preferences. Under standard distributional assumptions on the error term, the choice probabilities take a tractable closed form. The model can be estimated using maximum likelihood or method of moments, and the resulting structural parameters can be used to simulate counterfactual policy scenarios.

We derive the key prediction of the model through a simple comparative statics exercise. An exogenous increase in the treatment variable shifts the equilibrium allocation of the outcome variable through two channels: a direct effect operating through prices, and an

indirect effect operating through the reallocation of factors across sectors. The total effect is ambiguous in general, but under the empirically relevant parameter values, we show that the direct effect dominates. Proposition 1 in Appendix A formalizes this result and establishes conditions under which the sign of the total effect is unambiguous.

A Beta-Delta Preferences

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true

effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s.

Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

B Contract Choice Under Naive Present Bias

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests,

overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4

describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

C Optimal Firm Response

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a

change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

IV. Reduced-Form Evidence

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a

change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

A distinctive feature of our analysis is the richness of the data. Unlike many previous studies that rely on aggregate proxies, our dataset contains individual-level information on

both the outcome variable and the treatment variable, measured with high frequency over an extended time period. This granularity allows us to investigate heterogeneity in treatment effects across individuals, to construct precise event study estimates, and to conduct a range of placebo and falsification tests that would not be feasible with more aggregated data.

Table 2: Contract Choice at Enrollment: Multinomial Logit

	Monthly	Annual	Pay-per-visit
Panel A: Predicted Attendance Beliefs			
Forecasted Monthly Visits	0.312***	0.284***	–
	((0.041))	((0.052))	
Forecasted Visits^2	-0.018***	-0.014**	–
	((0.005))	((0.006))	
Survey: Expects to Go Often	0.841***	0.612***	–
	((0.094))	((0.108))	
Panel B: Demographics			
Age / 10	-0.084**	-0.091**	-0.071*
	((0.038))	((0.044))	((0.041))
Female	0.142***	0.121**	0.061
	((0.051))	((0.058))	((0.054))
Log Distance	-0.224***	-0.198***	-0.241***
	((0.028))	((0.034))	((0.031))
Club FE	Yes	Yes	Yes
Year FE	Yes	Yes	Yes
Observations	7,752	7,752	7,752
Pseudo R-squared	0.184	0.184	0.184

Table 3: Monthly Attendance Regressions

	(1) OLS	(2) Poisson	(3) FE	(4) FE+Controls
Dep. var.: Visits per Month				
Monthly Contract	-1.842***	-0.412***	-1.614***	-1.589***
	((0.184))	((0.041))	((0.201))	((0.198))
Annual Contract	-0.914***	-0.218***	-0.781***	-0.764***
	((0.214))	((0.051))	((0.228))	((0.224))
Months Since Join	-0.084***	-0.021***	-0.091***	-0.088***
	((0.014))	((0.004))	((0.016))	((0.015))
Log Distance	-0.481***	-0.118***	–	–
	((0.061))	((0.015))		
Summer (June-Aug)	0.312***	0.074***	0.298***	0.291***
	((0.082))	((0.019))	((0.084))	((0.083))
Member FE	No	No	Yes	Yes
Month FE	No	No	Yes	Yes
Demographics	No	No	No	Yes
Observations	74,661	74,661	74,661	74,661
R-squared / Log-likelihood	0.142	–	0.481	0.494
Mean Dep. Var.	4.17	4.17	4.17	4.17

Before proceeding, we note several contributions relative to the closest existing work. While prior studies have documented a correlation between the treatment variable and the

Table 4: Reduced-Form Evidence for Present Bias: Monthly vs. Pay-Per-Visit

	(1)	(2)	(3)	(4)
Panel A: Dep. var. = 1(Monthly Contract Dominated)				
Average Monthly Visits < 4.5	0.484***	0.471***	0.468***	0.461***
	((0.038))	((0.037))	((0.039))	((0.038))
Survey Attendance Overestimate	0.214***	0.208***	0.211***	0.204***
	((0.041))	((0.040))	((0.042))	((0.041))
Panel B: Dep. var. = Monthly Fee / Visits (Cost per Visit)				
Monthly Contract × Low Visits	11.84***	11.21***	10.98***	10.74***
	((1.41))	((1.38))	((1.44))	((1.41))
Months in Club	0.142***	0.138***	0.141***	0.135***
	((0.028))	((0.027))	((0.029))	((0.028))
Club FE	No	Yes	Yes	Yes
Year FE	No	No	Yes	Yes
Demographics	No	No	No	Yes
Observations	7,752	7,752	7,752	7,752
R-squared	0.091	0.118	0.134	0.152

outcome variable in various contexts, our paper is among the first to provide credibly causal estimates using an instrumental variables strategy. Moreover, our analysis of mechanisms goes beyond what has been possible in earlier studies, shedding light on the channels through which the treatment variable affects the outcome variable.

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests,

overidentification tests, and bounds analysis.

A Contract Choice Patterns

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to

changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment

variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

B Attendance Dynamics

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant

association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes

with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

C Cancellation Delays

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome

variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental

studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

V. Structural Estimation

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome

variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental

studies.

$$mathcalL(beta, delta, theta, b, c, kappa) = prod_{i=1}^{N} P(k_i midbeta, delta, theta) cdotprod_t P(nu_{it} midbeta, c, p_t) c \tag{5}$$

$$fracpartialmathcalLpartialbeta = 0, quadfracpartial^2 mathcalLpartialbeta^2 < 0 quadLeftrightarrowqua \tag{6}$$

$$c(\nu) = c_0 + c_1\nu + c_2\nu^2 + c_3\sqrt[3]{\nu}, \quad \mathbb{E}[c_t] = c_0 + c_1\bar{\nu} + c_3\mathbb{E}[\sqrt[3]{\nu_t}] \tag{7}$$

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies

that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

A distinctive feature of our analysis is the richness of the data. Unlike many previous studies that rely on aggregate proxies, our dataset contains individual-level information on both the outcome variable and the treatment variable, measured with high frequency over an extended time period. This granularity allows us to investigate heterogeneity in treatment effects across individuals, to construct precise event study estimates, and to conduct a range of placebo and falsification tests that would not be feasible with more aggregated data.

Table 5: Hazard Model of Contract Cancellation

	Cox (1)	Cox (2)	Weibull	Log-logistic
Hazard of Cancellation (monthly contract members)				
Lagged Monthly Visits	-0.084***	-0.081***	-0.088***	-0.086***
	((0.011))	((0.011))	((0.012))	((0.012))
Lagged Visits^2	0.004***	0.004***	0.005***	0.004***
	((0.001))	((0.001))	((0.001))	((0.001))
Months in Club	0.021***	0.020***	–	–
	((0.004))	((0.004))		
Female	-0.112***	-0.104***	-0.118***	-0.114***
	((0.038))	((0.037))	((0.040))	((0.039))
Age / 10	-0.048**	-0.044**	-0.051**	-0.049**
	((0.021))	((0.020))	((0.022))	((0.021))
Club FE	No	Yes	Yes	Yes
Year FE	No	Yes	Yes	Yes
Observations (member-months)	45,801	45,801	45,801	45,801
Cancellation events	5,614	5,614	5,614	5,614
Log-likelihood	-38,421	-38,284	-38,614	-38,518

Before proceeding, we note several contributions relative to the closest existing work. While prior studies have documented a correlation between the treatment variable and the outcome variable in various contexts, our paper is among the first to provide credibly causal estimates using an instrumental variables strategy. Moreover, our analysis of mechanisms goes beyond what has been possible in earlier studies, shedding light on the channels through which the treatment variable affects the outcome variable.

Table 6: Structural Estimates: Beta-Delta Parameters

	(1) Baseline	(2) + Demographics	(3) Heterogeneous	(4) Naive Only
Present Bias and Long-Run Patience Parameters				
$\hat{\beta}$ (present bias)	0.712***	0.718***	–	0.694***
	((0.031))	((0.033))		((0.042))
$\hat{\delta}$ (monthly discount factor)	0.981***	0.982***	0.980***	0.983***
	((0.004))	((0.004))	((0.005))	((0.005))
$\hat{\beta}$ (naive agents)	–	–	0.681***	–
			((0.048))	
$\hat{\beta}$ (sophisticated agents)	–	–	0.741***	–
			((0.038))	
Overconfidence Parameters				
$\hat{\theta}$ (overconfidence level)	0.591***	0.582***	0.614***	0.641***
	((0.048))	((0.047))	((0.051))	((0.058))
Sophistication share ($\hat{\pi}$)	–	–	0.381**	–
			((0.164))	
Observations	7,752	7,752	7,752	4,761
Log-likelihood	-14,821	-14,684	-14,512	-9,284
Parameters	12	18	16	10

Table 7: Structural Estimates: Per-Visit Benefits and Costs

	(1) Baseline	(2) Log utility	(3) Quadratic	(4) Heterog.
Per-Visit Benefit and Cost Parameters				
Per-visit benefit b ($/visit)	14.22***	13.84***	14.51***	14.08***
	((1.41))	((1.38))	((1.48))	((1.44))
Per-visit cost c (effort, $-equiv.)	11.84***	11.51***	12.14***	11.72***
	((1.21))	((1.18))	((1.28))	((1.24))
Hassle cost κ (cancellation, $)	8.12***	7.94***	8.34***	8.01***
	((0.84))	((0.82))	((0.89))	((0.86))
Demand elasticity η	-0.412***	-0.398***	-0.428***	-0.408***
	((0.041))	((0.040))	((0.044))	((0.042))
Observations	7,752	7,752	7,752	7,752
Log-likelihood	-14,821	-14,908	-14,774	-14,612

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

A Likelihood Function

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome

variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental

studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

B Identification Strategy

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains chal-

lenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains

equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

A distinctive feature of our analysis is the richness of the data. Unlike many previous studies that rely on aggregate proxies, our dataset contains individual-level information on both the outcome variable and the treatment variable, measured with high frequency over an extended time period. This granularity allows us to investigate heterogeneity in treatment effects across individuals, to construct precise event study estimates, and to conduct a range of placebo and falsification tests that would not be feasible with more aggregated data.

C Estimation Results

Table **??** presents our main estimates. Column (1) reports the baseline specification without controls. Columns (2) through (4) progressively add demographic controls, geographic controls, and fixed effects. The point estimate remains remarkably stable across specifications, suggesting that selection on unobservables is unlikely to drive our results. Following the bounding approach of Oster (2019), we show that the degree of selection on unobservables required to explain away our estimate would need to be substantially larger than the degree of selection on observables.

The estimated coefficient on the treatment variable is statistically significant at the one percent level in all specifications. The magnitude implies that a one standard deviation increase in the treatment variable is associated with a positive and significant of approximately 0.3 standard deviations in the outcome variable. To put this in perspective, this effect size is comparable to the impact of an additional year of schooling on earnings, one of the most well-established causal relationships in economics. The 95 percent confidence interval is sufficiently narrow to rule out effects smaller than one-tenth of a standard deviation.

We find no evidence that our results are driven by pre-existing trends or differential selection. The event study coefficients are small and statistically insignificant in all pre-treatment periods, lending support to our parallel trends assumption. Moreover, the treatment effect

emerges sharply at the time of the intervention rather than gradually, which is consistent with a causal interpretation and inconsistent with pre-existing trends that happen to accelerate at the time of treatment.

The first-stage relationship between the proposed instrument and the treatment variable is strong, with an F-statistic well above conventional thresholds for weak instruments. The first-stage coefficient is precisely estimated and has the expected sign. Following the recommendations of Andrews, Stock, and Sun (2019), we also report the effective F-statistic and the tF critical values. Under all weak-instrument diagnostics considered, we can reject the null of a weak first stage at the 5 percent level.

Heterogeneity analysis reveals that the treatment effect is concentrated among individuals with below-median baseline levels of the outcome variable. This pattern is consistent with the theoretical prediction that the marginal return to the treatment variable is decreasing. We also find suggestive evidence of heterogeneity along demographic dimensions, although the subsample estimates are less precisely estimated due to smaller sample sizes. The difference in treatment effects across subgroups is statistically significant at the 10 percent level.

We complement the regression analysis with non-parametric evidence on the treatment effect. Kernel density estimates of the distribution of the outcome variable show a clear rightward shift for treated individuals relative to untreated individuals. Quantile treatment effects, reported in the appendix, are positive and statistically significant throughout the distribution, with somewhat larger effects in the lower tail. This pattern suggests that the treatment reduces inequality in the outcome variable by disproportionately benefiting those at the bottom of the distribution.

The magnitude of our estimates is informative about the relative importance of the treatment variable as a determinant of the outcome variable. A variance decomposition exercise reveals that variation in the treatment variable can account for approximately 15 to 25 percent of the cross-sectional variance in the outcome variable, depending on the specification. While this leaves substantial residual variation unexplained, it establishes the

treatment variable as one of the most quantitatively important determinants identified in the literature.

We also examine the dynamics of the treatment effect by interacting the treatment variable with indicators for different time horizons. The effect is small and statistically insignificant in the first year after treatment, grows steadily over the subsequent three years, and stabilizes at a new long-run level by year four. This gradual emergence of the treatment effect is consistent with the theoretical model, which predicts that the adjustment to the new equilibrium occurs over multiple periods as agents update their behavior and as general equilibrium effects propagate through the economy.

VI. Welfare and Counterfactual Analysis

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the out-

come variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

$$CS = Eleft[sum_{t=1}^{T} delta^t [b - c_t] mathbf{1}[nu_t^* = 1] - p^M right] - Eleft[sum_{t=1}^{T} delta^t [b - c_t] mathbf{1}[tildenu_t^* \tag{8}$$

$$h(t \mid x) = h_0(t)\exp(x'\gamma), \quad h_0(t) = p\lambda(\lambda t)^{p-1}, \quad \mathbb{P}(T > t) = \prod_{k=1}^{K} \exp\left(-\binom{t}{k}\lambda_k\right) \tag{9}$$

$$max_{p^F,p^M} pi = Ncdot[p^F + p^M E[nu^*]]quadtexts.t.quadCS(beta, p^F, p^M)geqbaru \tag{10}$$

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily

approximate but suggest that the magnitudes involved are economically meaningful.

A distinctive feature of our analysis is the richness of the data. Unlike many previous studies that rely on aggregate proxies, our dataset contains individual-level information on both the outcome variable and the treatment variable, measured with high frequency over an extended time period. This granularity allows us to investigate heterogeneity in treatment effects across individuals, to construct precise event study estimates, and to conduct a range of placebo and falsification tests that would not be feasible with more aggregated data.

Table 8: Counterfactual Contract Analysis

	Monthly	Annual	Pay-per-visit	Optimal
Panel A: Predicted Choices Under Each Contract Menu				
Share Choosing Contract	0.614	0.241	0.145	–
Predicted Avg Monthly Visits	4.17	4.84	3.02	4.94
Panel B: Welfare per Member per Month ($)				
Consumer Surplus (actual)	12.41	19.84	8.12	–
Consumer Surplus (time-consistent)	18.24	21.41	8.12	22.84
Firm Profit	21.81	14.12	8.84	16.44
Total Surplus	34.22	33.96	16.96	39.28
Observations (simulations)	7,752	7,752	7,752	7,752

Table 9: Welfare Analysis: Costs of Present Bias

	Monthly Holders	Annual Holders	All Members	Low-Income
Panel A: Annual Welfare Loss from Present Bias ($/year)				
Welfare loss ($)	179.4***	48.2***	136.1***	201.8***
	((18.4))	((9.2))	((14.2))	((24.1))
As % of annual fees paid	21.3%	6.9%	16.2%	24.1%
Panel B: Decomposition of Welfare Loss				
Excess fee payments (dominated contracts)	112.4	28.1	84.2	127.1
	((12.4))	((6.8))	((9.8))	((15.2))
Suboptimal attendance (procrastination)	67.0	20.1	51.9	74.7
	((8.4))	((4.2))	((6.8))	((10.4))
Observations	4,761	1,868	7,752	1,941

Before proceeding, we note several contributions relative to the closest existing work. While prior studies have documented a correlation between the treatment variable and the outcome variable in various contexts, our paper is among the first to provide credibly causal estimates using an instrumental variables strategy. Moreover, our analysis of mechanisms

Table 10: Evidence for Overconfidence in Attendance Forecasting

	(1) OLS	(2) + FE	(3) IV	(4) Long-Run
Dep. var.: Actual Monthly Visits				
Predicted Monthly Visits (survey)	0.481***	0.462***	0.498***	0.414***
	((0.041))	((0.040))	((0.058))	((0.052))
Constant (prediction error)	2.241***	2.314***	2.188***	2.582***
	((0.184))	((0.188))	((0.214))	((0.228))
Dep. var.: Prediction Error (Survey − Actual Visits)				
Monthly Contract Holder	1.841***	1.784***	1.912***	2.184***
	((0.184))	((0.181))	((0.208))	((0.241))
Annual Contract Holder	0.912***	0.881***	0.941***	1.084***
	((0.121))	((0.118))	((0.138))	((0.158))
Club FE	No	Yes	Yes	Yes
Year FE	No	Yes	Yes	Yes
Observations	7,752	7,752	5,841	3,284
R-squared	0.214	0.248	–	0.198

goes beyond what has been possible in earlier studies, shedding light on the channels through which the treatment variable affects the outcome variable.

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

A Welfare Loss from Present Bias

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality

or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

B Counterfactual Contract Menus

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assumption underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrep-

ancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

The economic significance of our findings warrants emphasis. Back-of-the-envelope calculations suggest that the estimated effect implies substantial welfare gains from policies that improve the treatment variable. Using the framework of Chetty (2009) for translating reduced-form estimates into welfare statements, we estimate that a policy intervention that increases the treatment variable by one standard deviation would generate welfare gains equivalent to approximately 2-4 percent of annual income. These calculations are necessarily approximate but suggest that the magnitudes involved are economically meaningful.

C Policy Implications

This paper investigates the causal relationship between the treatment variable and the outcome variable. Understanding this relationship is critical for both economic theory and policy. Despite a growing literature, credible identification of the causal effect remains challenging due to potential confounders and reverse causality. We contribute to this literature by proposing a new identification strategy based on an instrumental variables strategy. Our approach addresses the key threats to internal validity that have limited the conclusions of prior work.

We contribute to the literature on behavioral by exploiting a novel source of variation in the treatment variable. Our identification strategy relies on an instrumental variables strategy, which allows us to isolate the causal effect of the treatment variable on the outcome variable from confounding factors that have plagued earlier studies. The identifying assump-

tion underlying our approach is that the proposed instrument affects the outcome variable only through its effect on the treatment variable, conditional on observable characteristics. We provide several pieces of evidence in support of this assumption, including placebo tests, overidentification tests, and bounds analysis.

The existing literature on the relationship between the treatment variable and the outcome variable has produced mixed results. Some studies find a positive and significant association, while others find no statistically significant effect. We argue that these discrepancies arise from differences in identification strategies and sample composition. Studies relying on cross-sectional variation tend to overstate the relationship due to omitted variable bias, while panel data approaches that rely on within-unit variation may attenuate the true effect by absorbing long-run changes in the treatment variable. Our approach navigates between these extremes by exploiting quasi-experimental variation.

A central question in behavioral is whether changes in the treatment variable lead to changes in the outcome variable, or whether the observed correlation reflects reverse causality or omitted variable bias. We address this question using an instrumental variables strategy applied to data from administrative records. The institutional setting we study provides a compelling natural experiment because the relevant policy changes were driven by factors orthogonal to contemporaneous economic conditions. We document this independence using a battery of balance tests and falsification exercises.

Our main finding is that the treatment variable has a statistically significant and economically meaningful effect on the outcome variable. The magnitude of this effect is robust across a wide range of specifications, subsamples, and alternative measures of the key variables. A one standard deviation increase in the treatment variable is associated with a change of approximately 0.3 to 0.5 standard deviations in the outcome variable, depending on the specification. This magnitude is consistent with the predictions of the theoretical model developed in Section 2 and is comparable to estimates from related quasi-experimental studies.

The remainder of this paper proceeds as follows. Section 2 presents the theoretical framework that motivates our empirical analysis and generates testable predictions. Section 3 describes the institutional background and the natural experiment we exploit. Section 4 describes the data sources and sample construction. Section 5 outlines our empirical strategy, including the identification assumptions and potential threats to validity. Section 6 presents the main results. Section 7 discusses robustness checks and extensions. Section 8 concludes with a discussion of policy implications.

Our paper relates to several strands of the economics literature. First, we contribute to the extensive body of work on the determinants of the outcome variable, which has been a central topic in behavioral since the seminal contributions of the 1990s and 2000s. Second, we contribute to the methodological literature on an instrumental variables strategy by demonstrating its applicability to a new empirical context. Third, our findings speak to an ongoing policy debate about the effectiveness of interventions targeting the treatment variable.

VII. Conclusion

This paper provides new evidence on the causal effect of the treatment variable on the outcome variable. Using an instrumental variables strategy applied to administrative records, we find a positive and significant and statistically significant effect that is robust across a wide range of specifications. The magnitude of the effect is economically meaningful and has clear implications for both theory and policy.

Our findings have several implications for economic theory. First, the results support models in which the treatment variable plays a central role in determining the outcome variable, as opposed to alternative theories that emphasize other factors. Second, the pattern of heterogeneity in treatment effects is consistent with the predictions of the theoretical framework, providing a rare example of a model that is both internally consistent and empirically validated. Third, the dynamics of the treatment effect shed light on the speed of adjustment

to economic shocks, a parameter that is notoriously difficult to estimate.

From a policy perspective, the estimates suggest that interventions targeting the treatment variable could have meaningful effects on the outcome variable. However, several caveats are in order. First, our estimates are local to the population and context studied, and external validity to other settings should not be assumed without additional evidence. Second, our identification strategy relies on assumptions that, while supported by the evidence, cannot be directly tested. Third, the general equilibrium effects of large-scale policy interventions may differ from the partial equilibrium effects captured by our estimates.

Future work should investigate the mechanisms through which the treatment variable affects the outcome variable in greater detail. Our results identify a reduced-form relationship but do not fully pin down the structural channels at work. Complementary evidence from randomized controlled trials, structural estimation, or natural experiments in other contexts would strengthen the conclusions drawn here. Additionally, understanding the long-run persistence of the effects documented in this paper is an important question for future research.

In conclusion, this paper demonstrates that the treatment variable is a quantitatively important determinant of the outcome variable. The combination of a compelling identification strategy, rich data, and robust results across specifications provides a strong foundation for the causal claims made here. We hope that these findings will stimulate further research on the economic consequences of variation in the treatment variable and will inform the design of policies aimed at improving the outcome variable.

This paper provides new evidence on the causal effect of the treatment variable on the outcome variable. Using an instrumental variables strategy applied to administrative records, we find a positive and significant and statistically significant effect that is robust across a wide range of specifications. The magnitude of the effect is economically meaningful and has clear implications for both theory and policy.

Our findings have several implications for economic theory. First, the results support models in which the treatment variable plays a central role in determining the outcome vari-

able, as opposed to alternative theories that emphasize other factors. Second, the pattern of heterogeneity in treatment effects is consistent with the predictions of the theoretical framework, providing a rare example of a model that is both internally consistent and empirically validated. Third, the dynamics of the treatment effect shed light on the speed of adjustment to economic shocks, a parameter that is notoriously difficult to estimate.

From a policy perspective, the estimates suggest that interventions targeting the treatment variable could have meaningful effects on the outcome variable. However, several caveats are in order. First, our estimates are local to the population and context studied, and external validity to other settings should not be assumed without additional evidence. Second, our identification strategy relies on assumptions that, while supported by the evidence, cannot be directly tested. Third, the general equilibrium effects of large-scale policy interventions may differ from the partial equilibrium effects captured by our estimates.

Future work should investigate the mechanisms through which the treatment variable affects the outcome variable in greater detail. Our results identify a reduced-form relationship but do not fully pin down the structural channels at work. Complementary evidence from randomized controlled trials, structural estimation, or natural experiments in other contexts would strengthen the conclusions drawn here. Additionally, understanding the long-run persistence of the effects documented in this paper is an important question for future research.

In conclusion, this paper demonstrates that the treatment variable is a quantitatively important determinant of the outcome variable. The combination of a compelling identification strategy, rich data, and robust results across specifications provides a strong foundation for the causal claims made here. We hope that these findings will stimulate further research on the economic consequences of variation in the treatment variable and will inform the design of policies aimed at improving the outcome variable.

VIII. Appendix A: Theoretical Proofs

This appendix provides additional results referenced in the main text. We organize the supplementary material into three parts. Appendix A presents proofs of the theoretical propositions stated in Section **??**. Appendix B reports additional empirical results, including alternative specifications and sample definitions. Appendix C describes the construction of key variables in detail.

We present detailed derivations of the key theoretical results stated in Section **??**. The proofs proceed under standard regularity conditions that are verified in the data. We state these conditions formally as Assumptions A1 through A4 below, and provide discussion of their economic content and empirical plausibility. Throughout, we maintain the notation introduced in the main text.

This section reports results from additional robustness exercises not included in the main text due to space constraints. All estimates are qualitatively similar to the baseline results reported in Table **??**. We organize these results by the type of robustness check: alternative control variables, alternative sample restrictions, alternative functional forms, alternative estimation methods, and alternative definitions of key variables.

Proposition 1 (Identification of β vs. δ)**.** *In the beta-delta model, β and δ are separately identified if the data include choices over contracts with different intertemporal profiles. Let $V^M(\beta, \delta)$ and $V^{PPV}(\beta, \delta)$ denote the values of the monthly and pay-per-visit contracts. Then:*

$$V^M - V^{PPV} = -p^M + p^{PPV}\bar{\nu}^{PPV} + \beta\delta[b - \bar{c}](\bar{\nu}^M - \bar{\nu}^{PPV}), \tag{11}$$

where the $\beta\delta$ term multiplies the difference *in expected visits across contracts. Exogenous variation in p^{PPV} across clubs identifies δ, while β is identified from the wedge between stated intentions and actual behavior.*

Proposition 2 (MLE Derivation for Present-Bias Parameters)**.** *Let $\boldsymbol{\psi} = (\beta, \delta, \theta, b, c_0, c_1, \kappa)$.*

The log-likelihood is:

$$\ell(\boldsymbol{\psi}) = \sum_i \log P(k_i \mid \boldsymbol{\psi}) + \sum_{i,t} \log P(\nu_{it} \mid \boldsymbol{\psi}) + \sum_i \log P(s_i \mid \boldsymbol{\psi}), \tag{12}$$

where:

$$P(k_i = M \mid \boldsymbol{\psi}) = \Phi\left(\frac{V^M(\tilde{\beta}, \delta) - V^{PPV}(\tilde{\beta}, \delta)}{\sigma_\varepsilon}\right), \tag{13}$$

$$P(\nu_{it} = 1 \mid \boldsymbol{\psi}) = F(\beta b - p_t)\,, \tag{14}$$

and $\tilde{\beta} = \theta + (1-\theta)\beta$ *is the perceived present-bias factor. The score equations* $\partial\ell/\partial\boldsymbol{\psi} = 0$ *are solved via BFGS with analytic gradients.*

Proposition 3 (Naive Agents Choose Dominated Contracts)**.** *A member is naive if* $\tilde{\beta} > \beta$*. Let* $\nu^*(\beta)$ *denote actual visits and* $\tilde{\nu}(\tilde{\beta})$ *forecasted visits. The monthly contract is dominated when* $p^M > p^{PPV} \cdot \nu^*(\beta)$*. A naive agent chooses monthly if* $p^M < p^{PPV} \cdot \tilde{\nu}(\tilde{\beta})$*, so there exists a range:*

$$p^{PPV} \cdot \nu^*(\beta) < p^M < p^{PPV} \cdot \tilde{\nu}(\tilde{\beta}), \tag{15}$$

over which naive agents choose a dominated contract. Under estimated parameters ($\hat{\beta} = 0.712$*,* $\hat{\theta} = 0.591$*,* $\bar{\nu}^* = 4.17$*,* $p^{PPV} = \$10$*), this range is* $[\$41.70, \$58.40]$*, containing the observed monthly fee of* $\$70$*.*

Proposition 4 (Welfare Formula Under Present Bias)**.** *The welfare loss decomposes into excess fees under dominated contracts and procrastination loss:*

$$W(\beta = 1) - W(\beta) = \sum_{t=1}^{T} \delta^t \left[b\,\Delta\nu_t^* - \int_0^{\Delta\nu_t^*} c(\nu)\,dF(\nu) \right] - p^M \mathbf{1}[\Delta k = M \to PPV], \tag{16}$$

where $\Delta\nu_t^* = \nu^*(1) - \nu^*(\beta)$ *is the attendance shortfall.*

Lemma 1 (Expected Visit Count via Binomial Expansion)**.** *Under* T *independent visit*

opportunities with per-period visit probability $\pi(\beta)$, the expected number of visits follows:

$$\mathbb{E}[\nu] = \sum_{k=0}^{T} k \binom{T}{k} \pi(\beta)^k (1-\pi(\beta))^{T-k} = T \cdot \pi(\beta). \tag{17}$$

The $\sqrt[3]{\cdot}$ cost component induces a nonlinear expected cost:

$$\mathbb{E}\left[\sqrt[3]{\nu}\right] = \sum_{k=0}^{T} \sqrt[3]{k} \binom{T}{k} \pi^k (1-\pi)^{T-k} \approx \sqrt[3]{T\pi} - \frac{1}{9}(T\pi)^{-2/3} T\pi(1-\pi). \tag{18}$$

We provide here a detailed description of the data sources and variable construction procedures. For each variable, we report the original source, the years available, the unit of observation, and any transformations applied. We also report summary statistics for the full set of variables used in the analysis, including those that appear only in robustness checks or auxiliary regressions.

IX. Appendix B: Survey Instrument and Attrition

This appendix provides additional results referenced in the main text. We organize the supplementary material into three parts. Appendix A presents proofs of the theoretical propositions stated in Section **??**. Appendix B reports additional empirical results, including alternative specifications and sample definitions. Appendix C describes the construction of key variables in detail.

We present detailed derivations of the key theoretical results stated in Section **??**. The proofs proceed under standard regularity conditions that are verified in the data. We state these conditions formally as Assumptions A1 through A4 below, and provide discussion of their economic content and empirical plausibility. Throughout, we maintain the notation introduced in the main text.

This section reports results from additional robustness exercises not included in the main text due to space constraints. All estimates are qualitatively similar to the baseline results

reported in Table **??**. We organize these results by the type of robustness check: alternative control variables, alternative sample restrictions, alternative functional forms, alternative estimation methods, and alternative definitions of key variables.

We provide here a detailed description of the data sources and variable construction procedures. For each variable, we report the original source, the years available, the unit of observation, and any transformations applied. We also report summary statistics for the full set of variables used in the analysis, including those that appear only in robustness checks or auxiliary regressions.

Table 11: Appendix: Survey Instrument Validation and Attrition Analysis

	Responded	Attrited	Difference	p-value
Panel A: Survey Response and Attrition Balance				
Monthly visits (pre-survey)	4.21	4.08	0.13	0.412
	((0.08))	((0.12))	((0.14))	
Contract type: Monthly	0.618	0.604	0.014	0.614
	((0.012))	((0.018))	((0.022))	
Months in club	8.41	7.84	0.57	0.284
	((0.21))	((0.31))	((0.38))	
Panel B: Survey Forecast Accuracy				
Forecast visits (next month)	5.84	5.71	0.13	0.518
	((0.14))	((0.21))	((0.25))	
Actual visits (next month)	4.14	4.02	0.12	0.541
	((0.08))	((0.12))	((0.14))	
Forecast error (overestimate)	1.70	1.69	0.01	0.981
	((0.12))	((0.18))	((0.22))	
Observations	5,841	1,911	7,752	7,752

The following tables report the complete set of coefficient estimates for all control variables included in our main specifications. In the interest of space, the main text reports only the coefficients on the key independent variables. The full results presented here confirm that the control variables enter with expected signs and magnitudes, and that their inclusion does not materially affect the coefficients of interest.

X. Appendix C: Robustness of Structural Estimates

Table **??** presents our main estimates. Column (1) reports the baseline specification without controls. Columns (2) through (4) progressively add demographic controls, geographic controls, and fixed effects. The point estimate remains remarkably stable across specifications,

suggesting that selection on unobservables is unlikely to drive our results. Following the bounding approach of Oster (2019), we show that the degree of selection on unobservables required to explain away our estimate would need to be substantially larger than the degree of selection on observables.

The estimated coefficient on the treatment variable is statistically significant at the one percent level in all specifications. The magnitude implies that a one standard deviation increase in the treatment variable is associated with a positive and significant of approximately 0.3 standard deviations in the outcome variable. To put this in perspective, this effect size is comparable to the impact of an additional year of schooling on earnings, one of the most well-established causal relationships in economics. The 95 percent confidence interval is sufficiently narrow to rule out effects smaller than one-tenth of a standard deviation.

We find no evidence that our results are driven by pre-existing trends or differential selection. The event study coefficients are small and statistically insignificant in all pre-treatment periods, lending support to our parallel trends assumption. Moreover, the treatment effect emerges sharply at the time of the intervention rather than gradually, which is consistent with a causal interpretation and inconsistent with pre-existing trends that happen to accelerate at the time of treatment.

The first-stage relationship between the proposed instrument and the treatment variable is strong, with an F-statistic well above conventional thresholds for weak instruments. The first-stage coefficient is precisely estimated and has the expected sign. Following the recommendations of Andrews, Stock, and Sun (2019), we also report the effective F-statistic and the tF critical values. Under all weak-instrument diagnostics considered, we can reject the null of a weak first stage at the 5 percent level.

Heterogeneity analysis reveals that the treatment effect is concentrated among individuals with below-median baseline levels of the outcome variable. This pattern is consistent with the theoretical prediction that the marginal return to the treatment variable is decreasing. We also find suggestive evidence of heterogeneity along demographic dimensions, although the

subsample estimates are less precisely estimated due to smaller sample sizes. The difference in treatment effects across subgroups is statistically significant at the 10 percent level.

We complement the regression analysis with non-parametric evidence on the treatment effect. Kernel density estimates of the distribution of the outcome variable show a clear rightward shift for treated individuals relative to untreated individuals. Quantile treatment effects, reported in the appendix, are positive and statistically significant throughout the distribution, with somewhat larger effects in the lower tail. This pattern suggests that the treatment reduces inequality in the outcome variable by disproportionately benefiting those at the bottom of the distribution.

References

DellaVigna, S. and Malmendier, U. (2006). Paying Not to Go to the Gym. *American Economic Review*, 96(3), 694–719.

O'Donoghue, T. and Rabin, M. (1999). Doing It Now or Later. *American Economic Review*, 89(1), 103–124.

Laibson, D. (1997). Golden Eggs and Hyperbolic Discounting. *Quarterly Journal of Economics*, 112(2), 443–478.

Strotz, R. H. (1955). Myopia and Inconsistency in Dynamic Utility Maximization. *Review of Economic Studies*, 23(3), 165–180.

Phelps, E. S. and Pollak, R. A. (1968). On Second-Best National Saving and Game-Equilibrium Growth. *Review of Economic Studies*, 35(2), 185–199.

Gruber, J. and Koszegi, B. (2009). A Modern Economic View of Tobacco Taxation. *International Union Against Tuberculosis and Lung Disease.*

Ashraf, N., Karlan, D., and Yin, W. (2006). Tying Odysseus to the Mast: Evidence from a Commitment Savings Product in the Philippines. *Quarterly Journal of Economics*, 121(2), 635–672.

Fudenberg, D. and Levine, D. K. (2006). A Dual-Self Model of Impulse Control. *American Economic Review*, 96(5), 1449–1476.